

This is a companion notebook for the book [Deep Learning with Python, Second Edition](#). For readability, it only contains runnable code blocks and section titles, and omits everything else in the book: text paragraphs, figures, and pseudocode.

If you want to be able to follow what's going on, I recommend reading the notebook side by side with your copy of the book.

This notebook was generated for TensorFlow 2.6.

▼ The mathematical building blocks of neural networks

▼ A first look at a neural network

Loading the MNIST dataset in Keras

```
from tensorflow.keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

[+ Code](#)[+ Text](#)

```
train_images.shape
```

```
len(train_labels)
```

```
train_labels
```

```
test_images.shape
```

```
len(test_labels)
```

```
test_labels
```

The network architecture

```
from tensorflow import keras
from tensorflow.keras import layers
model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])
```

The compilation step

```
model.compile(optimizer="rmsprop",  
              loss="sparse_categorical_crossentropy",  
              metrics=["accuracy"])
```

Preparing the image data

```
train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype("float32") / 255  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype("float32") / 255
```

"Fitting" the model

```
model.fit(train_images, train_labels, epochs=5, batch_size=128)
```

Using the model to make predictions

```
test_digits = test_images[0:10]  
predictions = model.predict(test_digits)  
predictions[0]
```

```
predictions[0].argmax()
```

```
predictions[0][7]
```

```
test_labels[0]
```

Evaluating the model on new data

```
test_loss, test_acc = model.evaluate(test_images, test_labels)  
print(f"test_acc: {test_acc}")
```

▼ Data representations for neural networks

▼ Scalars (rank-0 tensors)

```
import numpy as np  
x = np.array(12)  
x
```

```
x.ndim
```

▼ Vectors (rank-1 tensors)

```
x = np.array([12, 3, 6, 14, 7])  
x
```

```
x.ndim
```

▼ Matrices (rank-2 tensors)

```
x = np.array([[5, 78, 2, 34, 0],  
              [6, 79, 3, 35, 1],  
              [7, 80, 4, 36, 2]])  
x.ndim
```

▼ Rank-3 and higher-rank tensors

```
x = np.array([[[5, 78, 2, 34, 0],  
               [6, 79, 3, 35, 1],  
               [7, 80, 4, 36, 2]],  
             [[5, 78, 2, 34, 0],  
               [6, 79, 3, 35, 1],  
               [7, 80, 4, 36, 2]],  
             [[5, 78, 2, 34, 0],  
               [6, 79, 3, 35, 1],  
               [7, 80, 4, 36, 2]]])  
x.ndim
```

▼ Key attributes

```
from tensorflow.keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
```

```
train_images.ndim
```

```
train_images.shape
```

```
train_images.dtype
```

Displaying the fourth digit

```
import matplotlib.pyplot as plt
digit = train_images[4]
plt.imshow(digit, cmap=plt.cm.binary)
plt.show()
```

```
train_labels[4]
```

▼ Manipulating tensors in NumPy

```
my_slice = train_images[10:100]
my_slice.shape
```

```
my_slice = train_images[10:100, :, :]
my_slice.shape
```

```
my_slice = train_images[10:100, 0:28, 0:28]
my_slice.shape
```

```
my_slice = train_images[:, 14:, 14:]
```

```
my_slice = train_images[:, 7:-7, 7:-7]
```

▼ The notion of data batches

```
batch = train_images[:128]
```

```
batch = train_images[128:256]
```

```
n = 3
batch = train_images[128 * n:128 * (n + 1)]
```

Real-world examples of data tensors

Vector data

Timeseries data or sequence data

Image data

Video data

▼ The gears of neural networks: tensor operations

▼ Element-wise operations

```
def naive_relu(x):
    assert len(x.shape) == 2
    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] = max(x[i, j], 0)
    return x

def naive_add(x, y):
    assert len(x.shape) == 2
    assert x.shape == y.shape
    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[i, j]
    return x

import time

x = np.random.random((20, 100))
y = np.random.random((20, 100))

t0 = time.time()
for _ in range(1000):
    z = x + y
    z = np.maximum(z, 0.)
print("Took: {:.2f} s".format(time.time() - t0))

t0 = time.time()
for _ in range(1000):
    z = naive_add(x, y)
    z = naive_relu(z)
print("Took: {:.2f} s".format(time.time() - t0))
```

▼ Broadcasting

```
import numpy as np
X = np.random.random((32, 10))
```

```

y = np.random.random((10,))

y = np.expand_dims(y, axis=0)

Y = np.concatenate([y] * 32, axis=0)

def naive_add_matrix_and_vector(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0]
    x = x.copy()
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            x[i, j] += y[j]
    return x

import numpy as np
x = np.random.random((64, 3, 32, 10))
y = np.random.random((32, 10))
z = np.maximum(x, y)

```

▼ Tensor product

```

x = np.random.random((32,))
y = np.random.random((32,))
z = np.dot(x, y)

def naive_vector_dot(x, y):
    assert len(x.shape) == 1
    assert len(y.shape) == 1
    assert x.shape[0] == y.shape[0]
    z = 0.
    for i in range(x.shape[0]):
        z += x[i] * y[i]
    return z

def naive_matrix_vector_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 1
    assert x.shape[1] == y.shape[0]
    z = np.zeros(x.shape[0])
    for i in range(x.shape[0]):
        for j in range(x.shape[1]):
            z[i] += x[i, j] * y[j]
    return z

def naive_matrix_vector_dot(x, y):
    z = np.zeros(x.shape[0])

```

```

for i in range(x.shape[0]):
    z[i] = naive_vector_dot(x[i, :], y)
return z

```

```

def naive_matrix_dot(x, y):
    assert len(x.shape) == 2
    assert len(y.shape) == 2
    assert x.shape[1] == y.shape[0]
    z = np.zeros((x.shape[0], y.shape[1]))
    for i in range(x.shape[0]):
        for j in range(y.shape[1]):
            row_x = x[i, :]
            column_y = y[:, j]
            z[i, j] = naive_vector_dot(row_x, column_y)
    return z

```

▼ Tensor reshaping

```
train_images = train_images.reshape((60000, 28 * 28))
```

```

x = np.array([[0., 1.],
              [2., 3.],
              [4., 5.]])

```

```
x.shape
```

```

x = x.reshape((6, 1))
x

```

```

x = np.zeros((300, 20))
x = np.transpose(x)
x.shape

```

Geometric interpretation of tensor operations

A geometric interpretation of deep learning

▼ The engine of neural networks: gradient-based optimization

What's a derivative?

Derivative of a tensor operation: the gradient

Stochastic gradient descent

▼ Chaining derivatives: The Backpropagation algorithm

The chain rule

Automatic differentiation with computation graphs

▼ The gradient tape in TensorFlow

```
import tensorflow as tf
x = tf.Variable(0.)
with tf.GradientTape() as tape:
    y = 2 * x + 3
grad_of_y_wrt_x = tape.gradient(y, x)
```

```
x = tf.Variable(tf.random.uniform((2, 2)))
with tf.GradientTape() as tape:
    y = 2 * x + 3
grad_of_y_wrt_x = tape.gradient(y, x)
```

```
W = tf.Variable(tf.random.uniform((2, 2)))
b = tf.Variable(tf.zeros((2,)))
x = tf.random.uniform((2, 2))
with tf.GradientTape() as tape:
    y = tf.matmul(x, W) + b
grad_of_y_wrt_W_and_b = tape.gradient(y, [W, b])
```

▼ Looking back at our first example

```
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images.reshape((60000, 28 * 28))
train_images = train_images.astype("float32") / 255
test_images = test_images.reshape((10000, 28 * 28))
test_images = test_images.astype("float32") / 255

model = keras.Sequential([
    layers.Dense(512, activation="relu"),
    layers.Dense(10, activation="softmax")
])
```



```

model.compile(optimizer="rmsprop",
              loss="sparse_categorical_crossentropy",
              metrics=["accuracy"])

model.fit(train_images, train_labels, epochs=5, batch_size=128)

```

▼ Reimplementing our first example from scratch in TensorFlow

▼ A simple Dense class

```

import tensorflow as tf

class NaiveDense:
    def __init__(self, input_size, output_size, activation):
        self.activation = activation

        w_shape = (input_size, output_size)
        w_initial_value = tf.random.uniform(w_shape, minval=0, maxval=1e-1)
        self.W = tf.Variable(w_initial_value)

        b_shape = (output_size,)
        b_initial_value = tf.zeros(b_shape)
        self.b = tf.Variable(b_initial_value)

    def __call__(self, inputs):
        return self.activation(tf.matmul(inputs, self.W) + self.b)

    @property
    def weights(self):
        return [self.W, self.b]

```

▼ A simple Sequential class

```

class NaiveSequential:
    def __init__(self, layers):
        self.layers = layers

    def __call__(self, inputs):
        x = inputs
        for layer in self.layers:
            x = layer(x)
        return x

    @property
    def weights(self):
        weights = []
        for layer in self.layers:
            weights += layer.weights
        return weights

```

```

model = NaiveSequential([
    NaiveDense(input_size=28 * 28, output_size=512, activation=tf.nn.relu),
    NaiveDense(input_size=512, output_size=10, activation=tf.nn.softmax)
])
assert len(model.weights) == 4

```

▼ A batch generator

```

import math

class BatchGenerator:
    def __init__(self, images, labels, batch_size=128):
        assert len(images) == len(labels)
        self.index = 0
        self.images = images
        self.labels = labels
        self.batch_size = batch_size
        self.num_batches = math.ceil(len(images) / batch_size)

    def next(self):
        images = self.images[self.index : self.index + self.batch_size]
        labels = self.labels[self.index : self.index + self.batch_size]
        self.index += self.batch_size
        return images, labels

```

▼ Running one training step

```

def one_training_step(model, images_batch, labels_batch):
    with tf.GradientTape() as tape:
        predictions = model(images_batch)
        per_sample_losses = tf.keras.losses.sparse_categorical_crossentropy(
            labels_batch, predictions)
        average_loss = tf.reduce_mean(per_sample_losses)
    gradients = tape.gradient(average_loss, model.weights)
    update_weights(gradients, model.weights)
    return average_loss

```

```

learning_rate = 1e-3

```

```

def update_weights(gradients, weights):
    for g, w in zip(gradients, weights):
        w.assign_sub(g * learning_rate)

```

```

from tensorflow.keras import optimizers

```

```

optimizer = optimizers.SGD(learning_rate=1e-3)

```

```
def update_weights(gradients, weights):  
    optimizer.apply_gradients(zip(gradients, weights))
```

▼ The full training loop

```
def fit(model, images, labels, epochs, batch_size=128):  
    for epoch_counter in range(epochs):  
        print(f"Epoch {epoch_counter}")  
        batch_generator = BatchGenerator(images, labels)  
        for batch_counter in range(batch_generator.num_batches):  
            images_batch, labels_batch = batch_generator.next()  
            loss = one_training_step(model, images_batch, labels_batch)  
            if batch_counter % 100 == 0:  
                print(f"loss at batch {batch_counter}: {loss:.2f}")  
  
from tensorflow.keras.datasets import mnist  
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()  
  
train_images = train_images.reshape((60000, 28 * 28))  
train_images = train_images.astype("float32") / 255  
test_images = test_images.reshape((10000, 28 * 28))  
test_images = test_images.astype("float32") / 255  
  
fit(model, train_images, train_labels, epochs=10, batch_size=128)
```

▼ Evaluating the model

```
predictions = model(test_images)  
predictions = predictions.numpy()  
predicted_labels = np.argmax(predictions, axis=1)  
matches = predicted_labels == test_labels  
print(f"accuracy: {matches.mean():.2f}")
```

Summary

